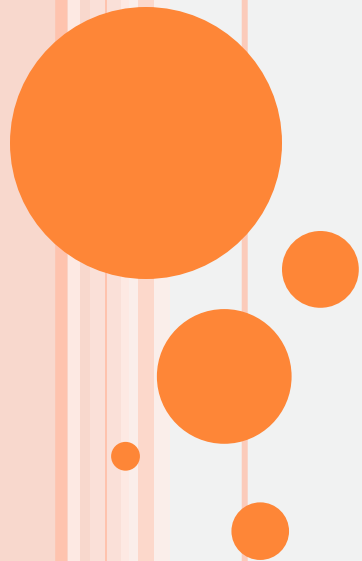


# CONCURRENCY CONTROL – 2PL



# DATABASE CONCURRENCY CONTROL

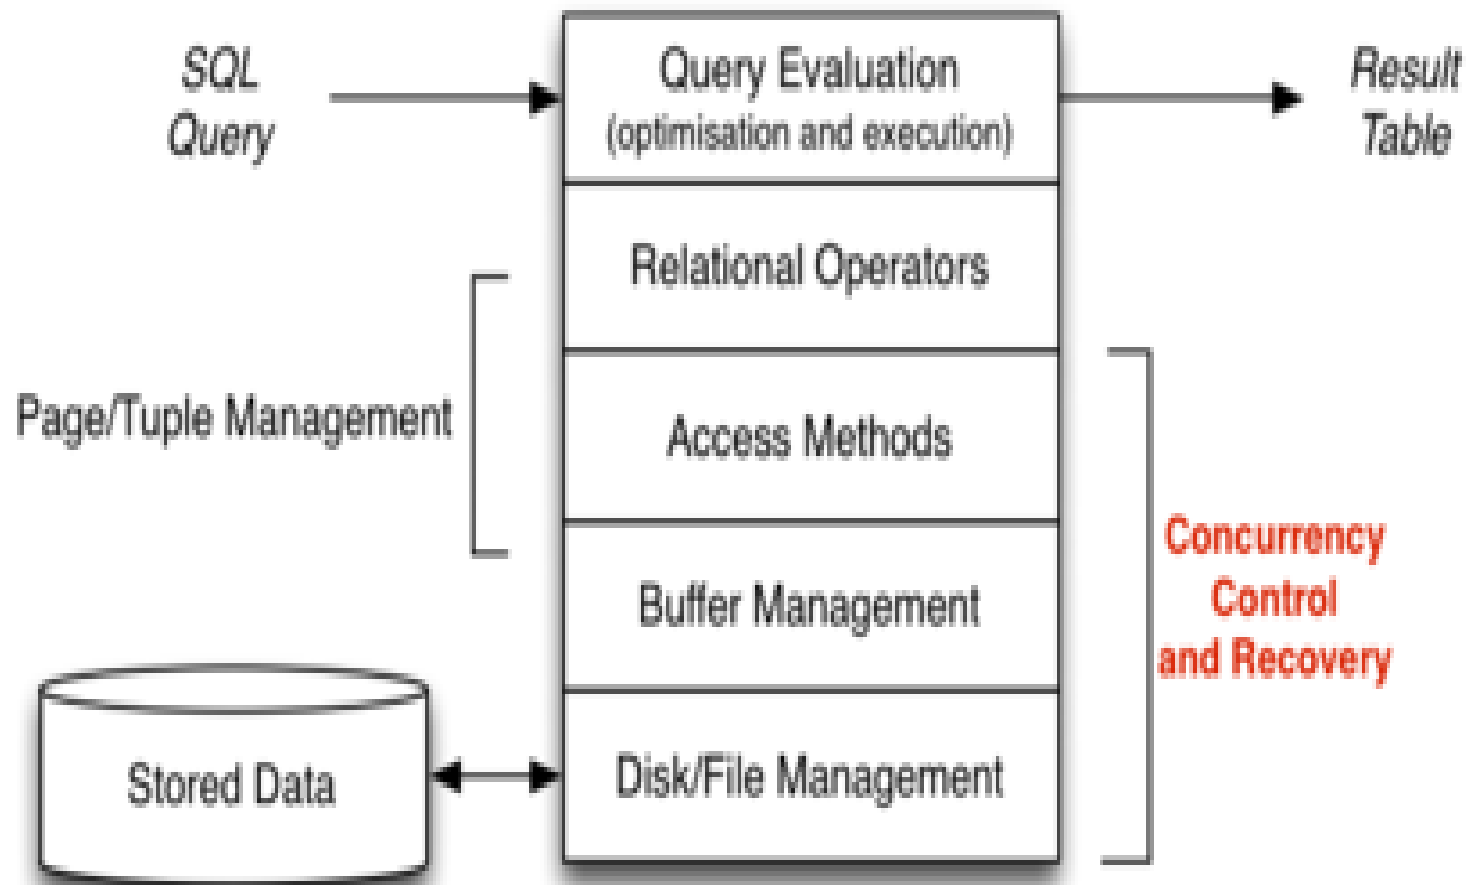
## Purpose of Concurrency Control

To enforce Isolation (through mutual exclusion) among conflicting transactions.

To preserve database consistency through consistency preserving execution of transactions.

To resolve read-write and write-write conflicts.





# TWO-PHASE LOCKING TECHNIQUE

Locking is an operation which secures

- (a) permission to Read
- (b) permission to Write a data item for a transaction.

**Example:** Lock (X).

Data item X is locked in behalf of the requesting transaction.

Unlocking is an operation which removes these permissions from the data item.

- **Example:** Unlock (X):  
Data item X is made available to all other transactions.

Lock and Unlock are Atomic operations.



# TWO-PHASE LOCKING TECHNIQUES: ESSENTIAL COMPONENTS



Two locks modes:

- (a) shared (read)      (b) exclusive (write).

Shared mode: shared lock (X)

- More than one transaction can apply share lock on X for reading its value but no write lock can be applied on X by any other transaction.

Exclusive mode: Write lock (X)

- Only one write lock on X can exist at any time and no shared lock can be applied by any other transaction on X.

Conflict matrix

|       | Read | Write |
|-------|------|-------|
| Read  | Y    | N     |
| Write | N    | N     |

# TWO-PHASE LOCKING TECHNIQUES: ESSENTIAL COMPONENTS



## Lock Manager:

- Managing locks on data items.

## Lock table:

- Lock manager uses it to store the identify of transaction locking a data item, the data item, lock mode and pointer to the next data item locked. One simple way to implement a lock table is through linked list.

| Transaction ID | Data item id | lock mode | Ptr to next data item |
|----------------|--------------|-----------|-----------------------|
| T1             | X1           | Read      | Next                  |



# TWO-PHASE LOCKING TECHNIQUES:



## Two Phases:

- (a) Locking (Growing)
- (b) Unlocking (Shrinking).

## Locking (Growing) Phase:

- A transaction applies locks (read or write) on desired data items one at a time.


## Unlocking (Shrinking) Phase:

- A transaction unlocks its locked data items one at a time.

## Requirement:

- For a transaction these two phases must be mutually exclusively, that is, during locking phase unlocking phase must not start and during unlocking phase locking phase must not begin.







# Two-Phase Locking Techniques: Example



## T1

```
read_lock (Y);  
read_item (Y);  
unlock (Y);  
write_lock (X);  
read_item (X);  
X:=X+Y;  
write_item (X);  
unlock (X);
```

## T2

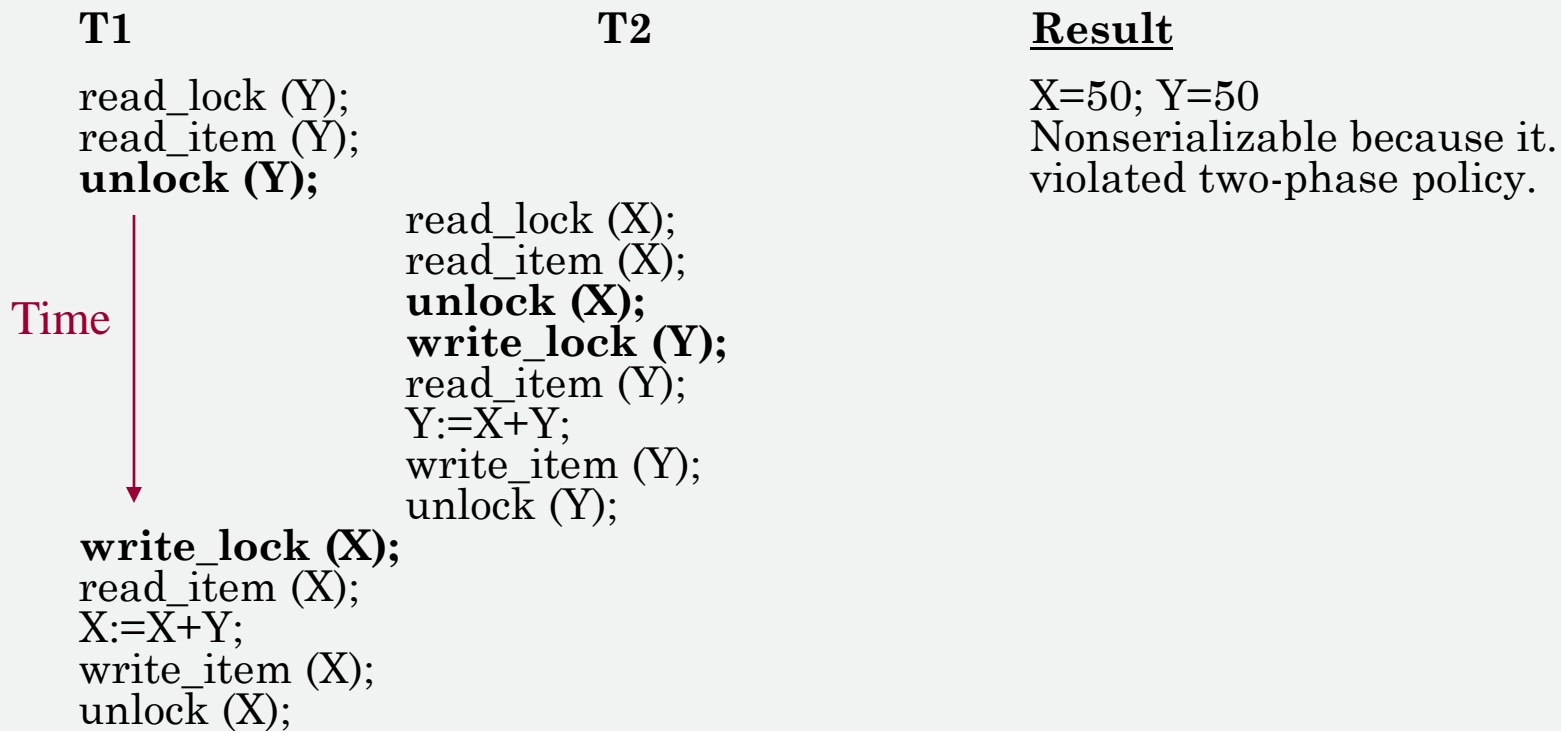
```
read_lock (X);  
read_item (X);  
unlock (X);  
Write_lock (Y);  
read_item (Y);  
Y:=X+Y;  
write_item (Y);  
unlock (Y);
```

## Result

Initial values: X=20; Y=30  
Result of serial execution  
T1 followed by T2  
X=50, Y=80.  
Result of serial execution  
T2 followed by T1  
X=70, Y=50



# Two-Phase Locking Techniques: Example



# TWO-PHASE LOCKING TECHNIQUES:

Two-phase policy generates two locking algorithms

- (a) **Basic**
- (b) **Conservative**

## Conservative:

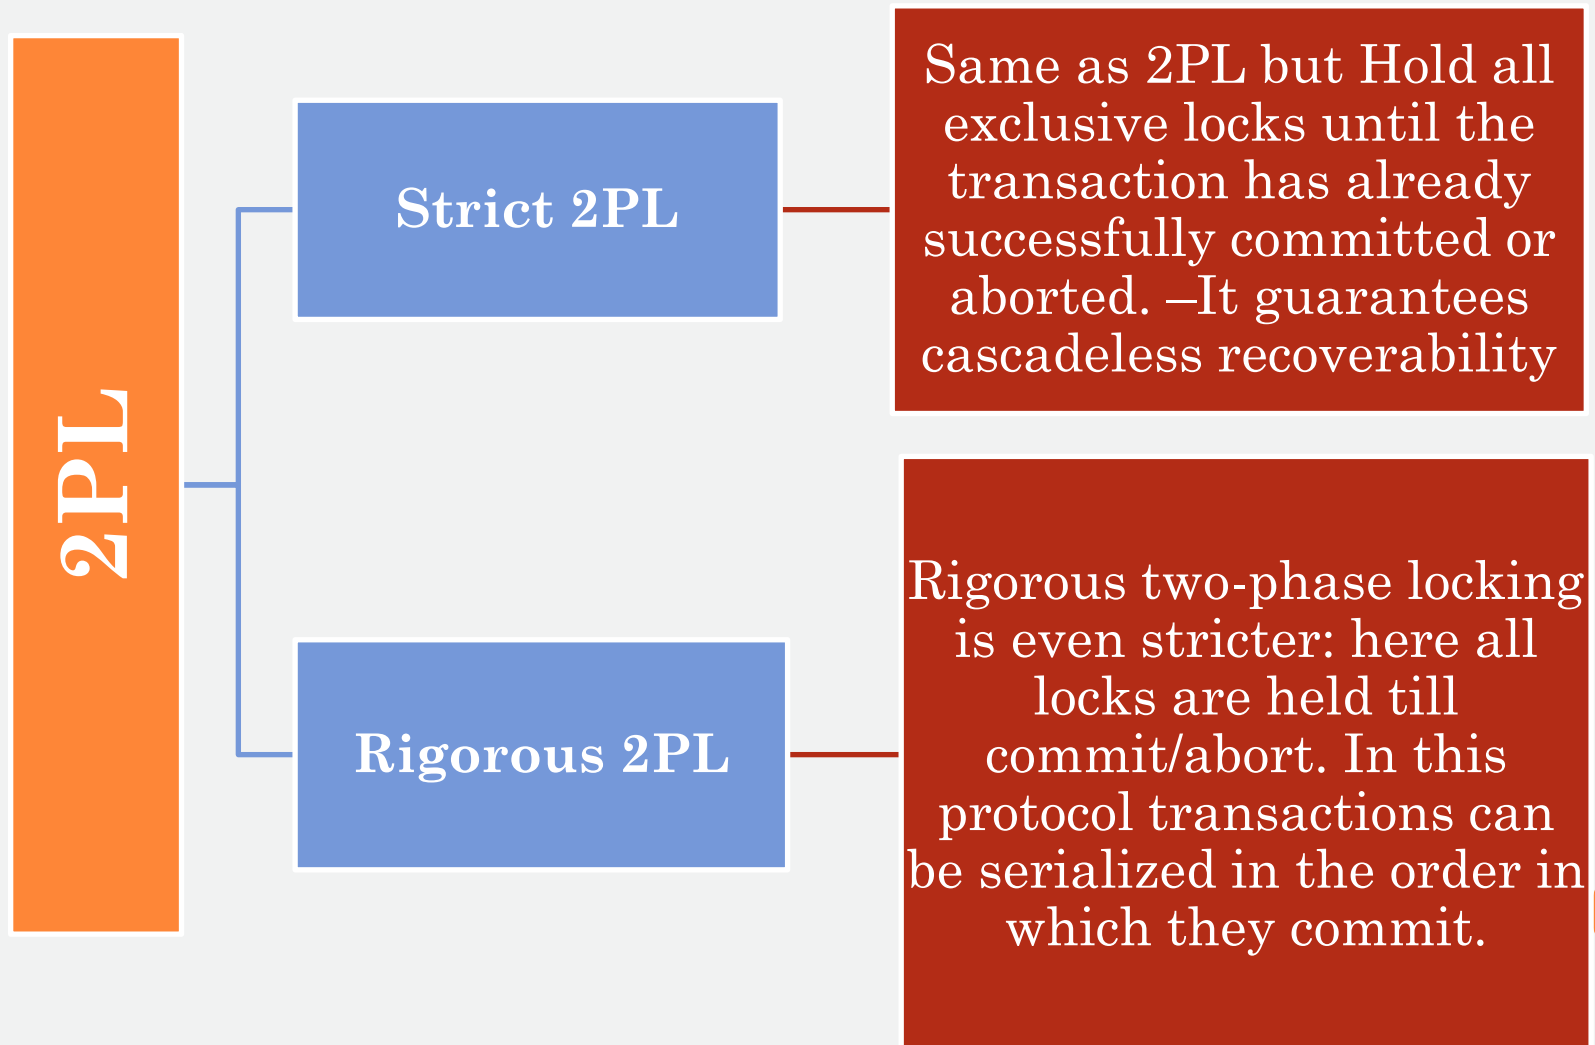
- Prevents deadlock by locking all desired data items before transaction begins execution.

## Basic:

- Transaction locks data items incrementally. This may cause deadlock which is dealt with.

## Strict:

- A more stricter version of Basic algorithm where unlocking is performed after a transaction terminates (commits or aborts and rolled-back). This is the most commonly used two-phase locking algorithm.



| Strict 2PL |           |
|------------|-----------|
| T1         | T2        |
| s-lock(A)  |           |
| read(A)    |           |
|            | s-lock(A) |
| x-lock(B)  |           |
| unlock(A)  |           |
| read(B)    |           |
| write(B)   |           |
|            | read(A)   |
|            | unlock(A) |
| commit     |           |
| unlock(B)  |           |
|            | s-lock(B) |
|            | read(B)   |
|            | unlock(B) |
|            | commit    |

| Rigorous 2PL |           |
|--------------|-----------|
| T1           | T2        |
| s-lock(A)    |           |
| read(A)      |           |
|              | s-lock(A) |
| x-lock(B)    |           |
|              | read(A)   |
| read(B)      |           |
| write(B)     |           |
| commit       |           |
| unlock(B)    |           |
|              | s-lock(B) |
|              | read(B)   |
| unlock(A)    |           |
|              | commit    |
|              | unlock(A) |
|              | unlock(B) |



# PRACTICE PROBLEM



Consider the following schedule involving two transactions  $T_1$  and  $T_2$ .

| $T_1$  | $T_2$  |
|--------|--------|
| R(A)   |        |
|        | R(A)   |
| W(A)   |        |
| commit |        |
|        | W(A)   |
|        | R(A)   |
|        | commit |

Identify the Type of Schedule?



# PRACTICE PROBLEM SOLUTION



Consider the following schedule involving two transactions  $T_1$  and  $T_2$ .

| $T_1$  | $T_2$  |
|--------|--------|
| R(A)   |        |
|        | R(A)   |
| W(A)   |        |
| commit |        |
|        | W(A)   |
|        | R(A)   |
|        | commit |

This is a strict schedule since  $T_2$  reads and writes A which is written by  $T_1$  only after the commit of  $T_1$ .



## PRACTICE PROBLEM

Consider the following database schedule with two transactions, T1 and T2.

**S = r2(X); r1(X); r2(Y); w1(X); r1(Y); w2(X); a1; a2;**  
where  $r_i(Z)$  denotes a read operation by transaction  $T_i$  on a variable  $Z$ ,  $w_i(Z)$  denotes a write operation by  $T_i$  on a variable  $Z$  and  $a_i$  denotes an abort by transaction  $T_i$ . Which one of the following statements about the above schedule is TRUE?

- (A) S is non-recoverable
- (B) S is recoverable, but has a cascading abort
- (C) S does not have a cascading abort
- (D) S is strict





# PRACTICE PROBLEM SOLUTION

| T1   | T2   |
|------|------|
|      | R(X) |
| R(X) |      |
|      | R(Y) |
| W(X) |      |
| R(Y) |      |
|      | W(X) |
| a1   |      |
|      | a2   |

T1 performs Write operation on X

T2 performs write operation on same variable X

# PRACTICE PROBLEM SOLUTION



As we can see in figure,

- T2 overwrites a value that T1 writes
- T1 aborts: its “remembered” values are restored.
- Cascading Abort could have arisen if – > Abort of T1 required abort of T2 but as T2 is already aborted, it's not a cascade abort. Therefore, **Option C**
- **Option A** – is **not** true because the given schedule is recoverable
- **Option B** – is **not** true as it is recoverable and avoids cascading aborts;
- **Option D** – is **not** true because T2 is also doing abort operation after T1 does, so NOT strict.

# PRACTICE PROBLEM



Consider the following schedule involving two transactions  $T_1$  and  $T_2$ .

Is this transaction implements 2PL?

|    | $T_1$     | $T_2$     |
|----|-----------|-----------|
| 1  | LOCK-S(A) |           |
| 2  |           | LOCK-S(A) |
| 3  | LOCK-X(B) |           |
| 4  | .....     | .....     |
| 5  | UNLOCK(A) |           |
| 6  |           | LOCK-X(C) |
| 7  | UNLOCK(B) |           |
| 8  |           | UNLOCK(A) |
| 9  |           | UNLOCK(C) |
| 10 | .....     | .....     |

# PRACTICE PROBLEM SOLUTION



Yes this transaction implements 2PL

This is just a skeleton transaction which shows how unlocking and locking works with 2-PL. Note for:

## **Transaction $T_1$ :**

Growing Phase is from steps 1-3.

Shrinking Phase is from steps 5-7.

Lock Point at 3

## **Transaction $T_2$ :**

Growing Phase is from steps 2-6.

Shrinking Phase is from steps 8-9.

Lock Point at 6



# DEALING WITH DEADLOCK AND STARVATION



- **Deadlock**

T'1

read\_lock (Y);  
read\_item (Y);

write\_lock (X);  
(waits for X)

T'2

read\_lock (X);  
read\_item (Y);

write\_lock (Y);  
(waits for Y)

T1 and T2 did follow two-phase policy but they are deadlock

- **Deadlock (T'1 and T'2)**



# DEALING WITH DEADLOCK AND STARVATION

## Deadlock prevention

A transaction locks all data items it refers to before it begins execution.

This way of locking prevents deadlock since a transaction never waits for a data item.

The conservative two-phase locking uses this approach.



# DEALING WITH DEADLOCK AND STARVATION

## Deadlock detection and resolution

In this approach, deadlocks are allowed to happen. The scheduler maintains a **wait-for-graph** for detecting cycle.

If a cycle exists, then one transaction involved in the cycle is selected (victim) and rolled-back.

A **wait-for-graph** is created using the lock table. As soon as a transaction is blocked, it is added to the graph. When a chain like:  $T_i$  waits for  $T_j$  waits for  $T_k$  waits for  $T_i$  or  $T_j$  occurs, then this creates a cycle.

# DEALING WITH DEADLOCK AND STARVATION

## Deadlock Avoidance

There are many variations of two-phase locking algorithm.

Some avoid deadlock by not letting the cycle to complete.

That is as soon as the algorithm discovers that blocking a transaction is likely to create a cycle, it rolls back the transaction.

Wound-Wait and Wait-Die algorithms use timestamps to avoid deadlocks by rolling-back victim.



# DEALING WITH DEADLOCK AND STARVATION

## Starvation

Starvation occurs when a particular transaction consistently waits or restarted and never gets a chance to proceed further.

In a deadlock resolution it is possible that the same transaction may consistently be selected as victim and rolled-back.

This limitation is inherent in all priority based scheduling mechanisms.

In Wound-Wait scheme a younger transaction may always be wounded (aborted) by a long running older transaction which may create starvation.

THANKS!!

